

Reinforcement **Learning** for Business, Economics, and Social Sciences

Unit 4-1: Neural Networks

Davud Rostam-Afschar (Uni Mannheim)

How to deal with very large
state-action spaces?

Tabular Value Iteration and Q-Learning

- ▶ Markov Decision Processes: value iteration

$$V(s) \leftarrow \max_a R(s) + \gamma \sum_{s'} \mathbb{P}(s' | s, a) V(s')$$


- ▶ Reinforcement Learning: Q-Learning

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

- ▶ Complexity depends on number of states and actions


Large State Spaces

- ▶ Computer Go: 3^{361} states



- ▶ Inverted pendulum: (x, x', θ, θ')

- ▶ 4-dimensional
- ▶ continuous state space



- ▶ Atari: $210 \times 160 \times 3$ dimensions (pixel values)



Functions to be Approximated

- ▶ Policy: $\pi(s) \rightarrow a$
- ▶ Value function: $V(s) \in \mathbb{R}$
- ▶ Q-function: $Q(s, a) \in \mathbb{R}$

Q-function Approximation

- ▶ Let $s = (x_1, x_2, \dots, x_n)$
→ states are defined by a vector of features x .
- ▶ Linear


$$Q(s, a) \approx \sum_i w_{ai} x_i$$

- ▶ Non-linear (e.g., neural network)

$$Q(s, a) \approx g(x; \mathbf{w})$$

Traditional Neural Network

- ▶ Network of units (computational neurons) linked by weighted edges



- ▶ Each unit computes:


$$z = h(\mathbf{w}'\mathbf{x} + b)$$

- ▶ Inputs: \mathbf{x}
- ▶ Output: z
- ▶ Weights (parameters): \mathbf{w}
- ▶ Bias: b
- ▶ Activation function (usually non-linear): h

Readings: Deep Neural Networks Goodfellow (2016, chapters 6, 7, 8)

One hidden Layer Architecture

- Feed-forward neural network



- Hidden units: $z_j = h_1 \left(\mathbf{w}_j^{'(1)} \mathbf{x} + b_j^{(1)} \right)$
- Output units: $y_k = h_2 \left(\mathbf{w}_k^{'(2)} \mathbf{z} + b_k^{(2)} \right)$
- Overall: $y_k = h_2 \left(\sum_j w_{kj}^{(2)} h_1 \left(\sum_i w_{ji}^{(1)} x_i + b_j^{(1)} \right) + b_k^{(2)} \right)$

Common Activation Functions

Common activation functions h

► Threshold:

$$h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

Common activation functions h

- Threshold:

$$h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

- Sigmoid:

$$h(a) = \sigma(a) = \frac{1}{1+e^{-a}}$$


Common activation functions h

► Threshold:

$$h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

► Sigmoid:

$$h(a) = \sigma(a) = \frac{1}{1+e^{-a}}$$



Common activation functions h

- ▶ Threshold:

$$h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

- ▶ Sigmoid:

$$h(a) = \sigma(a) = \frac{1}{1+e^{-a}}$$

- ▶ Gaussian:

$$h(a) = e^{-\frac{1}{2}\left(\frac{a-\mu}{\sigma}\right)^2}$$

Common activation functions h

- ▶ Threshold:

$$h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

- ▶ Sigmoid:

$$h(a) = \sigma(a) = \frac{1}{1+e^{-a}}$$

- ▶ Gaussian:

$$h(a) = e^{-\frac{1}{2}\left(\frac{a-\mu}{\sigma}\right)^2}$$

- ▶ Tanh : $h(a) =$

$$\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

Common activation functions h

- Threshold:

$$h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

- Sigmoid:


$$h(a) = \sigma(a) = \frac{1}{1+e^{-a}}$$

- Gaussian:

$$h(a) = e^{-\frac{1}{2}\left(\frac{a-\mu}{\sigma}\right)^2}$$

- Tanh : $h(a) =$

$$\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$



Common activation functions h

- ▶ Threshold:

$$h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

- ▶ Sigmoid:

$$h(a) = \sigma(a) = \frac{1}{1+e^{-a}}$$

- ▶ Gaussian:

$$h(a) = e^{-\frac{1}{2}\left(\frac{a-\mu}{\sigma}\right)^2}$$

- ▶ Tanh : $h(a) =$



$$\tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

- ▶ Identity: $h(a) = a$

Universal Function Approximation

Universal function approximation

- **Theorem:** Neural networks with at least one hidden layer of sufficiently many sigmoid/tanh/Gaussian units can approximate any function arbitrarily closely.



Minimize least squared error

- ▶ Minimize error function (Euclidian norm is commonly used for distance)

$$J(\mathbf{W}) = \frac{1}{2} \sum_n J_n(\mathbf{W})^2 = \frac{1}{2} \sum_n \|f(\mathbf{x}_n, \mathbf{W}) - y_n\|_2^2$$

where J is the error function, f is the function encoded by the neural net and n is the number data points.


- ▶ Train by gradient descent (a.k.a. backpropagation)
 - ▶ For each example (\mathbf{x}_n, y_n) , adjust the weights as follows:

$$w_{ji} \leftarrow w_{ji} - \alpha \frac{\partial J_n}{\partial w_{ji}}$$

α is the stepsize.

Minimize least squared error

- ▶ Train by gradient descent (a.k.a. backpropagation)
 - ▶ For each example (x_n, y_n) , adjust the weights as follows:



References I

GOODFELLOW, I. (2016): *Deep learning*, vol. 196. MIT press, Available at <http://deeplearningbook.org/>.

Takeaways

Neural Nets to Approximate Policies, Value or Quality Functions

- ▶ Tabular methods fail in large or continuous state-action spaces
- ▶ Neural networks approximate
 - ▶ policies,
 - ▶ value functions, and
 - ▶ Q-functions
- ▶ A basic network has
 - ▶ weighted inputs,
 - ▶ nonlinear activations, and
 - ▶ outputs
- ▶ Neural networks can approximate any continuous function (universal approximation)